

Benchmarking TCR–pMHC structure prediction without multiple sequence alignments: a protein language model matches AlphaFold 3, and their disagreement predicts accuracy

[Author list to be completed]

28 August 2026

Abstract

T cell receptor (TCR) recognition of peptide–MHC (pMHC) is the central specificity event of cellular immunity, and structure prediction is now the main route to modelling it at scale. Published benchmarks agree that multiple-sequence-alignment (MSA)-based predictors, and AlphaFold 3 in particular, lead the field, and that protein-language-model (PLM) predictors trail them. We revisited that conclusion on a benchmark of 136 TCR–pMHC entries released after the 30 September 2021 training cutoff shared by every method tested, of which 126 (111 class I, 15 class II) carry a scoreable experimental reference. All methods received byte-identical truncated inputs and a single fixed chain convention, and every model was scored with an explicit identity chain mapping against the same experimental references, cropped from the deposited coordinates to the input sequence and verified to draw all chains from a single author-defined assembly. AlphaFold 3 (median Global DockQ 0.751) and ESMFold2, a single-sequence diffusion predictor conditioned on a 6-billion-parameter language model rather than on an MSA (0.736), were statistically indistinguishable (paired Wilcoxon $p = 0.68$); both exceeded Protenix (0.701, $p = 0.045$ and $p = 0.0009$) and the TCR-specialised TCRmodel2 (0.708 on the 111 class I targets) was not separable from either. Equal accuracy did not mean equivalent behaviour. ESMFold2’s ipTM tracked accuracy far more tightly than AlphaFold 3’s (Pearson $r = 0.79$ vs 0.66); AlphaFold 3 was more sensitive to reference resolution; and the two models placed the TCR more than 90° apart on 10 of 126 complexes while agreeing on the pMHC to within 0.44 Å. The $C\alpha$ RMSD between the two predictions — a quantity computable without any experimental structure — predicted the worse of the two DockQ scores better than either model’s own confidence (Spearman $\rho = -0.884$ vs $+0.60$ and $+0.64$), and the two methods together reached medium-or-better accuracy on 119 of 126 complexes against 107 and 115 individually. Instrumenting both samplers showed the equal accuracy is reached by different routes: AlphaFold 3 is still 44 Å from its answer at the midpoint of denoising and converges only in the final fifth of its schedule, whereas ESMFold2 is within 6 Å at its midpoint. Exploiting that, cutting AlphaFold 3 from 200 to 160 denoising steps and from 5 diffusion samples to 1 gave a $1.80\times$ speed-up on 30 complexes at a cost of -0.010 Global DockQ (95% CI -0.032 to $+0.007$, $p = 0.33$). We conclude that MSAs are not required for state-of-the-art TCR–pMHC modelling, that cross-architecture agreement is a strong reference-free quality signal, and that the default AlphaFold 3 sampler is substantially over-provisioned for complexes of this size.

1 Introduction

The $\alpha\beta$ T cell receptor recognises a short peptide presented by a major histocompatibility complex (MHC) protein, and the geometry of that engagement — the docking angle of the TCR variable

domains over the peptide-binding platform, and the conformations of the six complementarity-determining region (CDR) loops that contact the peptide — determines specificity. Experimental determination of TCR–pMHC complexes remains slow relative to the size of the repertoire, so computational modelling has become the practical route to structural hypotheses at scale, whether for epitope-specificity prediction, TCR engineering, or immunotherapeutic design [1–3].

Deep-learning structure prediction transformed this problem twice. AlphaFold 2 made accurate single-chain and many multi-chain predictions routine [4], though early benchmarks identified adaptive immune recognition — antibody–antigen and TCR–pMHC complexes in particular — as a systematic failure mode [5]. AlphaFold 3 replaced the Structure Module with a diffusion head and reported substantially improved interface accuracy across biomolecular space [6], and TCR-specific adaptations such as TCRdock [7] and TCRmodel2 [8] narrowed the gap for this class of complex specifically. In parallel, protein language models made structure prediction possible without an MSA at all: ESMFold demonstrated that a sufficiently large language model encodes atomic-level structure in its representations [9], trading evolutionary coupling signal for a learned prior and an order-of-magnitude speed-up.

Recent TCR–pMHC benchmarks have converged on a consistent ranking. Lu et al. evaluated thirteen methods on 70 post-cutoff complexes and concluded that “MSA-based methods, particularly AlphaFold3, consistently outperform all alternatives”, with most PLM-based models underperforming except a PLM variant built on an AlphaFold-3-style architecture [10]. Asuncion-París et al. reached the same conclusion on 265 class I complexes, reporting AlphaFold 3 as the best generaliser to post-training structures [11]. Shi et al. found that TCR-specific tools derived from AlphaFold 2 gain interface accuracy at the cost of framework accuracy, and that CDR3 loops, docking orientation and class II MHC–peptide interfaces remain difficult for every method [12]. Yin et al. observed complementary behaviour between methods and pooling gains for antibody–peptide complexes [13]. Together these establish the field’s working assumption: evolutionary information is what makes TCR–pMHC prediction work.

Two considerations motivate re-examining that conclusion. First, the PLM predictors in those comparisons are predominantly ESMFold-generation models whose structure module is a direct AlphaFold-2 descendant, whereas a current-generation language-model predictor with a diffusion structure head is a substantially different architecture. Second, a benchmark that ranks methods addresses only part of the question. Where two architectures that share little on the conditioning side — one using a paired MSA and structural templates, the other only the query sequence — achieve comparable accuracy, the more informative questions concern the route each takes to that accuracy and whether the disagreements between them carry usable information.

Here we assemble a benchmark of 136 TCR–pMHC entries released after 30 September 2021, the training cutoff shared by AlphaFold 3, Protenix and ESMFold2, of which 126 carry a scoreable experimental reference. Every method receives identical truncated inputs under one chain convention, and every prediction is scored against a native cropped from the deposited coordinates and audited to come from a single author-defined assembly. We report interface accuracy (DockQ), TCR-specific accuracy (MHC-aligned weighted all-CDR C α RMSD and framework-aligned CDR3 all-atom RMSD), confidence calibration, sensitivity to reference resolution and to training-set similarity, the structural anatomy of the disagreement between the two leading models, the internal trajectories of both diffusion samplers, and the wall-clock cost of the default AlphaFold 3 sampler settings.

2 Materials and methods

2.1 Benchmark set construction

We constructed a dataset of TCR-pMHC Class I and Class II complexes from the TCR3d database [1,3], which curates all known TCR complex structures. Restricting to $\alpha\beta$ TCR complexes with an initial RCSB release date after 30 September 2021 gave 427 class I and 72 class II candidates. Of these, 136 were complete TCR-peptide-MHC complexes containing all of an MHC peptide-binding platform, a peptide, and paired TCR α and β variable domains (114 class I, 22 class II); the remainder were unliganded TCRs, pMHC-only entries, TCR-mimic antibody complexes, or otherwise incomplete triads.

All 136 entries were reviewed against their deposited coordinates and primary publications, and ten were excluded from scoring. Five (8TRL, 8TRQ, 8TRR, 9NIG, 9WNT) contain modified peptide residues that a standard FASTA sequence cannot represent, and because the modification is in every case part of the interaction being modelled, predicting these complexes from sequence alone is not well posed. Four (8VD0, 8VQ8, 9AUD, 9C3E) are engineered constructs in which the peptide is covalently fused to the MHC through a linker; modelling their components as separate chains would alter the connectivity of the experiment. The last, 9J4S, has no physical TCR-pMHC interface in any author-defined assembly of the deposited structure and so cannot be scored at all. This leaves 126 scoreable complexes, 111 class I and 15 class II.

No redundancy filter or resolution cut was applied to the primary set. TCR variable-domain identity is a poor novelty axis — germline V-gene framework dominates a ~ 110 -residue V domain — and filtering on it discards most of the set while retaining an idiosyncratic remainder (Section 3.10). Instead, similarity to the pre-cutoff structural record is measured explicitly and treated as a covariate.


2.2 Domain truncation and chain convention

Each complex was truncated to the domains that constitute the variable regions of the TCR-pMHC interface (Figure 1). For class I, $\beta 2$ -microglobulin was removed and the MHC heavy chain was cut to its $\alpha 1\alpha 2$ peptide-binding platform (median 181 residues, range 161–181); for class II, the $\alpha 1$ and $\beta 1$ domains were retained (81–85 and 89–93 residues). TCR chains were cut to their variable domains (α 101–114 residues, β 105–116). Prediction inputs used the complete deposited polymer sequences over those spans.


Every method received the same chain letters, with no per-model remapping anywhere in the pipeline: for class I, MHC $\alpha 1\alpha 2$ (A), peptide (B), TCR α (C) and TCR β (D); for class II, MHC-II α (A), peptide (B), MHC-II β (C), TCR α (D) and TCR β (E).

A Class I (PDB 7RTR)

Full experimental complex




Truncated benchmark input



B Class II (PDB 8PJG)

Full experimental complex



Truncated benchmark input




Figure 1: **Domain truncation applied to the benchmark complexes.** (A) Class I, PDB 7RTR. The TCR constant domains, the MHC $\alpha 3$ domain and $\beta 2$ -microglobulin are removed, leaving the MHC-I $\alpha 1\alpha 2$ peptide-binding platform, the peptide and the paired TCR variable domains. (B) Class II, PDB 8PJG. The TCR constant domains and the MHC-II $\alpha 2$ and $\beta 2$ domains are removed, leaving the $\alpha 1$ and $\beta 1$ peptide-binding domains, the peptide and the paired TCR variable domains. Retained chains are coloured by role; removed regions are grey.

2.3 Structure prediction methods

Four methods were run on identical inputs.

AlphaFold 3 [6] predictions were obtained through the AlphaFold Server, which supplies MSAs and structural templates from its own pipeline and returns five ranked diffusion samples; the top-ranked model by ranking score was scored. The server enforces a 30 September 2021 training cutoff.

Protenix [22], an open reproduction of the AlphaFold 3 architecture, was run through its public server with the same inputs, and the top-ranked model scored.

ESMFold2 (the `biohub/ESMFold2` checkpoint, conditioned on the `biohub/ESMC-6B` language model) [21] was run through the ESM SDK API as model `esmfold2-fast-2026-05` with the SDK defaults of 10 refinement loops and 100 sampling steps. ESMFold2 is a diffusion structure predictor in the same sense as AlphaFold 3, but it takes no MSA and no structural template: its only input is the query sequence, embedded by a 6-billion-parameter protein language model. This is the architectural contrast on which the present comparison rests, and it is a strict one, since the two methods share the diffusion generative head and little else on the conditioning side.

TCRmodel2 [8] was installed locally from the Pierce laboratory repository (AlphaFold 2.3 base, `alphafold_params_2022-12-06`) with `--max_template_date=2021-09-30` applied explicitly to both the TCR and the pMHC template searches; the default is effectively unrestricted, so this setting is required for a valid post-cutoff comparison. TCRmodel2 accepts a fixed peptide length for class II and was therefore run on class I only, giving 111 predictions.

Mean end-to-end runtimes were 5.2 s per complex for the ESMFold2 API (median, IQR 4.9–6.4 s) and 13.4 s per complex for a local AlphaFold 3 run at default sampler settings on one RTX PRO 6000 (Section 2.11).

2.4 Interface accuracy

Interface accuracy was measured with DockQ v2.1.3 [14,15], run once per complex. The primary reported quantity is DockQ v2’s Global DockQ, its interface-size weighted score over the whole complex; per-pairwise-interface DockQ values are also retained, giving 6 interfaces per class I complex and 10 per class II complex. Standard CAPRI thresholds are used throughout (< 0.23 incorrect, 0.23–0.49 acceptable, 0.49–0.80 medium, ≥ 0.80 high) [16].

2.5 CDR annotation

TCR α and TCR β chains were annotated once per biological sequence with ANARCI [17] under IMGT numbering [18], and the annotation propagated to prediction and native coordinates through pairwise sequence alignment rather than through PDB residue numbering. CDR ranges were the inclusive IMGT intervals 27–38 (CDR1), 56–65 (CDR2), 81–86 (CDR2.5) and 104–118 (CDR3); “all CDRs” means all four regions on both chains, eight in total.

2.6 TCR-specific RMSD metrics

Two TCR-specific metrics were computed on the same top-ranked predictions the DockQ workflow scores; predictions were never selected using a native RMSD. The first, the MHC-aligned weighted

all-CDR C α RMSD, follows the TCRdock formulation [7,19]. MHC residues are matched between prediction and native by sequence alignment; a single ordinary-least-squares Kabsch superposition of the class I peptide-binding platform (or, jointly, of both class II peptide-binding domains) is applied to the entire predicted complex, and no further fit is performed. C α RMSD is then computed over the eight CDR regions with CDR3 residues weighted 3 and all others weighted 1:

$$\text{RMSD}_w = \sqrt{\frac{\sum_i w_i d_i^2}{\sum_i w_i}}$$

Because there is no second fit, this metric measures how accurately the TCR is docked on the pMHC.

The second, the framework-aligned CDR3 all-atom RMSD, measures the opposite quantity. For each chain independently, the predicted variable-domain framework (the IMGT-numbered variable domain excluding all four CDR regions) is superimposed on the native framework using matched non-hydrogen atoms, and all-heavy-atom RMSD is then computed over the corresponding CDR3 loop without further alignment, comparing only identically named atoms in sequence-matched, identity-matched residues. C α -only versions are retained as quality-control columns.

2.7 Superposition-free CDR accuracy (IDDT)

Both metrics above answer their question through a superposition, and each superposition is a choice. The local distance difference test (IDDT) [23] removes that choice: a residue is scored by the fraction of its interatomic distances to the rest of the structure that the prediction reproduces, averaged over four tolerances (0.5, 1, 2 and 4 Å), counting only pairs closer than 15 Å in the reference and excluding pairs within one residue. It needs no alignment and no reference frame, and it is bounded in [0, 1].

IDDT was computed on C α atoms, which is AlphaFold’s own definition and requires no symmetric-side-chain correction; the all-atom view of the CDRs is already carried by the framework-aligned CDR3 RMSD. Residues were matched between prediction and native through the same manifest-sequence alignment, and the CDRs defined by the same ANARCI annotation, as the two RMSD metrics, so the residue sets behind the two families of number are identical. Every TCR region was scored twice, differing only in which atoms may act as contact partners: against every chain of the complex, which makes the score sensitive to where the TCR is docked, and against the CDR’s own TCR chain alone, which removes the docking and leaves local loop geometry. Whole-chain, peptide, MHC and whole-complex IDDT are reported alongside for context.

2.8 Metric validation and quality control

Every metric reports expected, matched and resolved residue counts, atom counts, coverage fraction and a pass/fail status; incomplete loops are kept separate from the complete-case analysis rather than silently scored. Validation tests confirm that a native scored against itself, and against a rigidly rotated and translated copy of itself, returns approximately zero; that displacing an intact TCR relative to a fixed MHC raises the MHC-aligned CDR RMSD while leaving framework-aligned CDR3 RMSDs near zero; that deforming a CDR3 loop raises its own framework-aligned RMSD; and that permuting chain identifiers changes nothing, because chain assignment is by sequence and

biological role. The lDDT implementation is checked against biotite’s implementation of Mariani et al. to within 10^{-6} , returns 1 for a structure scored against itself and against a rigidly rotated and translated copy of itself, and is unaffected by rigid motion of a group lying outside the 15 Å inclusion radius.

2.9 Confidence and reference-free quality scores

Each method’s own interface confidence (ipTM, and pTM where available) was taken from its output. Per-residue predicted lDDT (pLDDT) was read from the B-factor column of the same top-ranked prediction file that was scored, averaged over a residue’s non-hydrogen atoms and then over a region; all four methods write it on the 0–100 scale here, and anything reported in $[0, 1]$ is rescaled. Confidence is therefore compared with accuracy residue-for-residue over exactly the residue sets of Section 2.7, and the calibration error reported is the mean of $\text{pLDDT} - \text{lDDT}$ over complexes. pDockQ [20] was computed per interface from the predicted structure and per-residue confidence, and averaged per complex. A published random-forest quality-tier classifier for TCR–pMHC models [11] was re-implemented against this benchmark’s chain convention and scored on all 126 complexes; the subset of class I targets whose PDB identifiers do not appear in its training set ($n = 27$) is reported separately throughout.

2.10 A composite single-prediction confidence score

Each of the scores above summarises one aspect of a prediction. We combined four of them into a single quantity, MOGGER, as their geometric mean:

$$\text{MOGGER} = \left(\min_{i \neq j} \text{ipTM}_{ij} \cdot \text{pDockQ} \cdot \frac{\text{pLDDT}_{\text{CDR3}\alpha}}{100} \cdot \frac{1}{1 + \max_{i \neq j} \text{PAE}_{ij}^{\min}} \right)^{1/4}$$

where ipTM_{ij} and $\text{PAE}_{ij}^{\min}$ are the off-diagonal entries of AlphaFold 3’s chain-pair ipTM and chain-pair minimum-PAE matrices, so the first and fourth terms are the weakest and the least certain chain relationship in the complex; pDockQ is the interface score of Section 2.9; and $\text{pLDDT}_{\text{CDR3}\alpha}$ is the mean pLDDT of the CDR3 α loop, the region that has to be right for the TCR to dock correctly. Each term lies on the unit interval, so the geometric mean penalises a single weak term more than an arithmetic mean would — a prediction is only as good as its worst-supported part.

Every input is read from one AlphaFold 3 `model_0` prediction and its own summary and full confidence files. No second sample, no second architecture, no template and no reference structure enters the score, which is what makes it usable prospectively. The formula was fixed on this benchmark and is therefore reported as a benchmark result, not as a validated external predictor.

2.11 Sampler instrumentation and timing

Both models’ official samplers were patched observationally to retain the coordinate state after every production denoising step, changing no arithmetic: AlphaFold 3’s `diffusion_head.sample` already computes and discards the per-step positions inside its scan, and the patch keeps them;

ESMFold2’s `DiffusionStructureHead.sample` holds its state in a local, so its installed source was rewritten at a unique anchor. Parity between instrumented and stock runs under the same seed was verified per model (mean $|\Delta\text{DockQ}| = 0.0008$, maximum 0.0122).

Sampler economy was measured on 30 complexes drawn with a frozen seed from the 126 scoreable set (28 class I, 2 class II, 397–416 residues), run serially on one idle RTX PRO 6000 with identical MSAs and templates taken verbatim from the AlphaFold Server bundles the benchmark’s own predictions came from, and the same per-complex seed. Four tracks were run: $200 \text{ steps} \times 5 \text{ samples}$ (default), 160×5 , 200×1 and 160×1 . All 30 complexes fall in AlphaFold 3’s 512-token bucket, so JIT compilation and kernel autotuning happen once per track; both were charged to two warm-up predictions run and discarded before timing. Instrumentation costs $\sim 2\%$ and was run as a separate pass, so the baseline track is stock AlphaFold 3.

2.12 Statistics and software

Paired method comparisons use the Wilcoxon signed-rank test on the complexes both methods predicted. Correlations are Pearson and Spearman with bootstrap confidence intervals (10,000 resamples, fixed seed); regressions use ordinary least squares with HC3 robust standard errors.

3 Results

3.1 Composition of the benchmark set

The 126 scoreable complexes comprise 111 class I and 15 class II entries, 112 human and 14 mouse, released between 13 October 2021 and 29 July 2026. One hundred and twenty were determined by X-ray crystallography and six by cryo-EM, with resolutions from 1.54 to 4.44 Å (median 2.56 Å, IQR 2.20–3.00 Å). They span 27 distinct MHC alleles — the most frequent being HLA-A*02 (29), H2-Db (10), HLA-A*11 (10), HLA-B*27 (9), HLA-A*24 (9) and HLA-E (7) — and 79 distinct peptide sequences of 8–15 residues (median 9).

3.2 Overall interface accuracy

Every method produced at least an acceptable model of every complex: Global DockQ never fell below 0.23 for any method on any of the 126 complexes (Figure 2A). Median Global DockQ was 0.751 for AlphaFold 3 (IQR 0.571–0.826), 0.736 for ESMFold2 (0.592–0.840), 0.708 for TCRmodel2 on its 111 class I targets (0.590–0.808) and 0.701 for Protenix (0.530–0.799).

The two leading methods are not distinguishable. Paired over all 126 complexes, AlphaFold 3 minus ESMFold2 gives a mean difference of -0.016 and a median difference of -0.002 (Wilcoxon $p = 0.68$). Both, however, separate from Protenix: AlphaFold 3 $+0.039$ ($p = 0.045$) and ESMFold2 $+0.055$ ($p = 0.0009$). TCRmodel2 was not separable from any of the three on its class I subset (all $p > 0.23$).

The CAPRI-class view (Figure 2B) tells the same story with a different emphasis. No method produced an incorrect model. ESMFold2 placed the largest fraction in the medium-or-better band (91.3%: 54.0% medium, 37.3% high), followed by TCRmodel2 (88.3%: 60.4% medium, 27.9% high), AlphaFold 3 (84.9%: 50.0% medium, 34.9% high) and Protenix (79.4%: 54.0% medium, 25.4% high).

AlphaFold 3’s slightly higher median therefore coexists with a slightly heavier acceptable-only tail than ESMFold2’s — the two distributions cross.

This result runs counter to the current consensus [10,11]: a predictor that uses neither an MSA nor a structural template matches AlphaFold 3 on post-cutoff TCR–pMHC complexes.




Figure 2: Interface accuracy on 126 post-cutoff TCR–pMHC complexes. (A) Distribution of Global DockQ per method; box, IQR and median; dots, individual complexes; dashed lines, CAPRI thresholds. (B) The same data as CAPRI classes, as a percentage of each method’s own set. No method produced an incorrect model of any complex. TCRmodel2 covers the 111 class I complexes only, in both panels.

3.3 Accuracy by interface

Global DockQ is weighted across all interfaces of the complex, several of which are modelled accurately by every method. Decomposing it by interface (Table 1) localises the differences between methods. The MHC–peptide interface is effectively solved by every method: median DockQ 0.940 (AlphaFold 3), 0.939 (ESMFold2), 0.933 (Protenix) and 0.931 (TCRmodel2), with the class II MHC α –peptide and peptide–MHC β interfaces close behind (0.90–0.91). The TCR α –TCR β pairing is also reliable (0.79–0.82). Everything that separates methods, and everything that separates easy complexes from hard ones, is in the four TCR-versus-pMHC interfaces, whose medians run 0.55–0.71. One class II interface is a partial exception: AlphaFold 3 models the MHC-II α –MHC-II β heterodimer pairing itself at median DockQ 0.628, against 0.845 for ESMFold2 and 0.789 for Protenix, the single largest per-interface deficit for AlphaFold 3 in the table.

Table 1: **Median per-interface DockQ by method.** n is the number of complexes contributing that interface (an interface with no native contacts is not scored, so peptide-TCR interfaces are absent in a few complexes whose TCR does not touch the peptide). The final row collapses the four TCR-versus-pMHC interfaces to one value per complex, then takes the median over complexes.

Interface	n	AF3	Protenix	ESMFold2	TCRmodel2
MHC-I – peptide	111	0.940	0.933	0.939	0.931
MHC-II α – peptide	15	0.904	0.911	0.905	—
peptide – MHC-II β	15	0.909	0.887	0.902	—
TCR α – TCR β	126	0.798	0.800	0.821	0.791
MHC-II α – MHC-II β	15	0.628	0.789	0.845	—
MHC-I – TCR α	111	0.672	0.590	0.653	0.644
MHC-I – TCR β	110	0.686	0.551	0.696	0.611
peptide – TCR α	120	0.710	0.605	0.676	0.641
peptide – TCR β	122	0.687	0.633	0.690	0.681
MHC-II α – TCR α	11	0.615	0.600	0.714	—
MHC-II α – TCR β	12	0.773	0.767	0.839	—
MHC-II β – TCR α	15	0.720	0.687	0.819	—
MHC-II β – TCR β	15	0.719	0.519	0.764	—
All TCR-pMHC interfaces (per complex)	126	0.692	0.621	0.694	0.632

Collapsing those four interfaces to one value per complex, median TCR-versus-pMHC DockQ is 0.694 for ESMFold2, 0.692 for AlphaFold 3, 0.632 for TCRmodel2 and 0.621 for Protenix. On this stricter reading the acceptable-or-better rate falls from 100% to 92.9% (ESMFold2), 90.1% (TCRmodel2), 87.3% (AlphaFold 3) and 84.1% (Protenix), and the high-accuracy rate is 28.6%, 27.8%, 19.8% and 17.5% respectively. Approximately one complex in ten is therefore still docked incorrectly by the best-performing methods.

3.4 TCR-specific accuracy: docking versus CDR3 loop conformation

The two TCR-specific metrics decompose that residual error (Table 2).

Table 2: **TCR-specific structural accuracy, median (IQR) in ångströms, class I and II pooled.** The MHC-aligned metric applies one rigid superposition of the predicted MHC peptide-binding platform onto the native and then measures all eight CDRs with CDR3 residues weighted 3, so it reports docking accuracy. The framework-aligned metrics fit each TCR chain’s variable-domain framework independently and then measure the CDR3 loop, so they report loop accuracy independent of docking. TCRmodel2 is scored on its 111 class I predictions.

Metric	AF3	Protenix	ESMFold2	TCRmodel2
MHC-aligned weighted all-CDR C α RMSD	2.84	3.82	2.75	3.47
	(1.61-5.56)	(2.10-7.47)	(1.60-5.24)	(1.99-5.16)
Framework-aligned CDR3 α all-atom RMSD	1.94	2.33	1.81	2.02
	(1.35-2.88)	(1.53-3.08)	(1.25-2.59)	(1.47-3.08)
Framework-aligned CDR3 β all-atom RMSD	1.73	1.93	1.56	1.95
	(1.17-2.49)	(1.40-2.60)	(1.13-2.26)	(1.42-2.70)
n complexes	126	126	126	111

The MHC-aligned weighted all-CDR C α RMSD, which measures docking with no second fit, has a median of 2.75 Å for ESMFold2, 2.84 Å for AlphaFold 3 and 3.82 Å for Protenix over the 126 complexes, and 3.47 Å for TCRmodel2 over its 111 class I predictions. The distributions are strongly right-skewed — means are 5.45 (ESMFold2), 6.34 (AlphaFold 3), 7.72 (Protenix) and 5.59 Å (TCRmodel2), and every method’s maximum exceeds 50 Å — confirming that the aggregate is dominated by a minority of complexes in which the TCR is placed in an entirely wrong orientation, not by a uniform degradation.

Framework-aligned CDR3 RMSD, which removes docking error by construction, is far more uniform across methods. Median CDR3 α all-heavy-atom RMSD is 1.81 Å (ESMFold2), 1.94 Å (AlphaFold 3), 2.02 Å (TCRmodel2) and 2.33 Å (Protenix); median CDR3 β is 1.56 Å (ESMFold2), 1.73 Å (AlphaFold 3), 1.93 Å (Protenix) and 1.95 Å (TCRmodel2). The spread across methods for loop conformation (0.5 Å) is an order of magnitude smaller than the spread across complexes for docking (tens of ångströms), and the maxima are bounded — no method exceeds 7.9 Å on CDR3 α or CDR3 β . In other words, all four methods build approximately the same quality of CDR3 loop, and differ almost entirely in where they put it. This is consistent with Shi et al.’s observation that CDR3 loops and docking orientation are separate, and separately unsolved, problems [12].

Two secondary observations follow from Table 2. First, ESMFold2’s advantage is concentrated in the CDR3 β loop and in class II: its class II medians are 1.76 Å (all-CDR), 1.82 Å (CDR3 α) and 1.34 Å (CDR3 β), against AlphaFold 3’s 2.37, 2.37 and 2.05 Å. Second, restricting a model’s template search to TCR and MHC structures does not spare it the catastrophic failures. On the 111 class I complexes TCRmodel2’s all-CDR distribution is less dispersed than AlphaFold 3’s or Protenix’s (SD 7.59 Å against 9.31 and 9.34) but not than ESMFold2’s (7.03 Å), and its maximum is the same 52.7 Å as AlphaFold 3’s. Its advantage over the general models is a lower mean (5.59 Å against AlphaFold 3’s 6.18 on the same complexes), not an absence of wrong orientations.

3.5 Superposition-free CDR accuracy and predicted local confidence

Both metrics of Section 3.4 answer their question through a superposition, and each superposition is a choice. lDDT [23] removes that choice: it scores a residue by how many of its interatomic distances to the rest of the structure the prediction reproduces, so it needs no alignment and no reference frame. It is also the quantity every one of these methods predicts per residue and reports as pLDDT, which makes it the one accuracy measure directly comparable with the models’ own confidence. We computed C α lDDT over the same eight CDR regions, on the same top-ranked predictions, with two contact sets: the whole complex, and the CDR’s own TCR chain alone (Table 3, Figure 3).

Table 3: **CDR lDDT and pLDDT, median (IQR), class I and II pooled.** lDDT is superposition-free and bounded in $[0, 1]$; higher is better. “Whole complex” uses every chain as contact partners, so the score falls when the TCR is docked in the wrong place; “own chain” restricts the partners to the TCR chain the loop sits on, which removes the docking. pLDDT is each method’s own per-residue confidence over the same residues, rescaled to $[0, 1]$.

Region and contact set	AF3	Protenix	ESMFold2	TCRmodel2
All CDRs, complex	0.866 (0.793–0.907)	0.844 (0.784–0.891)	0.872 (0.819–0.911)	0.852 (0.806–0.893)
All CDRs, own chain	0.921 (0.884–0.948)	0.916 (0.883–0.939)	0.932 (0.894–0.956)	0.920 (0.886–0.943)
CDR3 α , complex	0.824 (0.714–0.899)	0.787 (0.720–0.870)	0.835 (0.733–0.894)	0.792 (0.735–0.881)
CDR3 β , complex	0.842 (0.773–0.918)	0.824 (0.745–0.898)	0.849 (0.783–0.934)	0.819 (0.769–0.890)
Framework, complex	0.961 (0.944–0.970)	0.958 (0.948–0.969)	0.960 (0.949–0.971)	0.961 (0.948–0.971)
Whole complex	0.935 (0.914–0.957)	0.931 (0.909–0.949)	0.940 (0.922–0.955)	0.942 (0.918–0.951)
Mean CDR3 pLDDT	0.874 (0.796–0.920)	0.890 (0.840–0.939)	0.784 (0.718–0.821)	0.826 (0.783–0.865)
<i>n</i> complexes	126	126	126	111

Within a method, all-CDR lDDT in the whole-complex frame is very nearly a restatement of docking quality: it correlates with Global DockQ at Pearson $r = 0.934$ (AlphaFold 3), 0.915 (ESMFold2), 0.907 (Protenix) and 0.904 (TCRmodel2), and with the MHC-aligned weighted all-CDR RMSD of Section 3.4 at Spearman $\rho = -0.911, -0.897, -0.887$ and -0.827 . Between methods it says much less. The four medians span 0.844 to 0.872 — a range of 0.028 — for CDRs whose MHC-aligned RMSD medians span 2.75 to 3.82 Å. lDDT is local and bounded, and a TCR placed in an entirely wrong orientation keeps most of the interatomic distances within its own chains, so the catastrophic minority that dominates the RMSD means is compressed into the lower whisker of Figure 3A rather than allowed to run to 50 Å. lDDT therefore ranks complexes well and separates methods poorly, and the two should not be substituted for one another.

Restricting the contact set to the CDR’s own TCR chain reproduces Section 3.4’s conclusion without any superposition at all. All-CDR lDDT rises to 0.916–0.932, a spread of 0.016 across four methods, and its correlation with Global DockQ falls from about 0.91 to 0.680 (AlphaFold 3), 0.674 (ESMFold2), 0.600 (TCRmodel2) and 0.557 (Protenix). The loops are built to nearly the same local standard everywhere; what differs is where they are put.

pLDDT ranks CDR3 accuracy better than any of the global confidence scores rank whole-complex accuracy. Pooling CDR3 α and CDR3 β , Pearson r between mean CDR3 pLDDT and CDR3 lDDT is 0.837 for ESMFold2, 0.750 for Protenix, 0.728 for TCRmodel2 and 0.724 for AlphaFold 3 — the same ordering, and for ESMFold2 a higher value, than the ipTM correlations of Section 3.8. Its

calibration, however, changes sign across the variable domain (Figure 3C). AlphaFold 3, Protenix and TCRmodel2 understate the framework (mean pLDDT – IDDT = -0.027 , -0.007 and $+0.004$) and overstate CDR3 ($+0.035$, $+0.084$ and $+0.011$); Protenix is the most overconfident loop-builder of the four, and the only method whose CDR3 pLDDT exceeds even its own-chain CDR3 IDDT. ESMFold2 is uniformly conservative — -0.100 on the framework rising to -0.058 on CDR3 — which is the per-residue counterpart of the ipTM intercept in Section 3.8: its confidence is on a lower scale but better ordered. Measured against the own-chain IDDT instead, every method is underconfident on CDR3 (-0.035 , $+0.005$, -0.125 and -0.060). Predicted IDDT therefore sits between the loop’s local quality and the quality it achieves once docked: the confidence head anticipates part of the docking error, but not the part that matters.




Figure 3: **Superposition-free CDR accuracy and the models’ own confidence.** (A) All-CDR C α IDDT with the whole complex as contact partners (dark) and with the CDR’s own TCR chain only (pale). Removing the docking raises every method by a similar amount and compresses the differences between them. (B) Mean CDR3 pLDDT against the CDR3 IDDT actually achieved, CDR3 α and CDR3 β pooled, with per-method linear fits and Pearson r ; the dashed line is perfect calibration. (C) Mean pLDDT – IDDT by region. Three of the four methods cross from understating the framework to overstating CDR3; ESMFold2 is conservative everywhere.

3.6 Effect of MHC class

Class II is not harder than class I on this benchmark (Figure 4C). Median Global DockQ for class II versus class I is 0.752 versus 0.750 (AlphaFold 3), 0.742 versus 0.736 (ESMFold2) and 0.713 versus 0.698 (Protenix); every difference is well inside the spread, and the class II subset is small ($n = 15$). ESMFold2’s class II mean (0.743) is its highest subgroup mean. We note this mainly as a caution: with 15 class II complexes, this benchmark cannot resolve a class effect of the size other studies report, and the TCRmodel2 comparison is class I only by construction.

3.7 Effect of reference resolution

Accuracy declines with worsening reference resolution for every method, but not equally (Figure 4B). Pearson r between resolution and Global DockQ is -0.292 for AlphaFold 3 ($p = 0.001$), -0.266 for ESMFold2 ($p = 0.003$), -0.238 for TCRmodel2 ($p = 0.012$) and -0.184 for Protenix ($p = 0.040$, Spearman $p = 0.065$, not significant). Part of this is unavoidable — a lower-resolution reference is a noisier target — but the ordering is informative: the two methods most sensitive to reference

quality are the two most accurate ones, which is what one expects if the residual error in the best models is approaching the coordinate uncertainty of the references themselves.

3.8 Confidence calibration and reference-free quality scores

Interface confidence predicts accuracy for all three general methods, but ESMFold2’s ipTM does so substantially better (Figure 4A). Pearson r between ipTM and Global DockQ is 0.793 for ESMFold2 (Spearman $\rho = 0.777$), against 0.660 for AlphaFold 3 ($\rho = 0.654$) and 0.657 for Protenix ($\rho = 0.679$). The regression lines also differ in intercept: ESMFold2 reaches a given DockQ at a systematically lower ipTM than the other two, so its confidence is both more conservative and better ordered. pDockQ shows the same ranking with lower correlations overall (ESMFold2 $r = 0.676$, Protenix 0.611, AlphaFold 3 0.577).




Figure 4: **Determinants of interface accuracy.** (A) Global DockQ against each method’s own ipTM, with per-method linear fits and Pearson r . ESMFold2’s confidence is both more conservative — it reaches a given DockQ at a lower ipTM — and better ordered. (B) Global DockQ against the resolution of the experimental reference. (C) Global DockQ by MHC class. Shaded bands are the CAPRI quality classes.

3.9 A composite single-prediction score beats the published classifier

For AlphaFold 3 specifically, single global confidence scores leave accuracy on the table. Replacing ipTM with the weakest *pairwise* ipTM already raises Pearson r from 0.660 to 0.695 over all 126 complexes, and the four-term composite of Section 2.10 — computed from one AlphaFold 3 prediction, with no second structure and no reference — reaches $r = 0.712$ (95% CI 0.617–0.791) and $\rho = 0.725$ (0.624–0.803) (Figure 5A).

It is the best of the scores available for an AlphaFold 3 prediction (Figure 5B). On the same 126 complexes the published random-forest quality-tier classifier [11] reaches $r = 0.662$ ($\rho = 0.674$), AlphaFold 3’s own ipTM 0.660 (0.654) and pDockQ 0.660 (0.679) — that is, the classifier does no better than the two confidence scores it was built to improve on. The paired difference between the composite and the classifier is +0.050 Pearson (95% bootstrap CI +0.015 to +0.085) and +0.051 Spearman (+0.008 to +0.098), so the advantage over all 126 complexes excludes zero on both metrics.

That margin narrows on every subset. On the 111 class I complexes the composite reaches $r = 0.738$ and $\rho = 0.760$ against the classifier’s 0.700 and 0.734, a paired +0.038 Pearson (CI +0.004 to +0.072) but +0.027 Spearman (−0.013 to +0.067). On the 27 class I targets whose PDB identifiers are absent from the classifier’s training set the two are indistinguishable, and the sign is not even consistent: +0.021 Pearson (−0.080 to +0.129) and −0.047 Spearman (−0.187 to +0.078). The 15 class II complexes favour the composite by +0.082 on both metrics, with intervals several times wider than the effect.

Two things follow. The composite is the better ranking of AlphaFold 3 predictions across this benchmark, but its formula was fixed on this benchmark, and 27 unseen targets cannot separate it from the classifier; we therefore do not claim a resolved advantage for any single-model score on unseen targets. And all of these scores — learned or not — are beaten by a quantity that needs no quality model at all (Section 3.11).




Figure 5: **A composite single-prediction score against the published alternatives.** (A) Global DockQ against MOGGER for the 126 complexes, with the linear fit and both correlations; shaded bands are the CAPRI quality classes. (B) Correlation with Global DockQ for every score computable from one AlphaFold 3 prediction, over all 126 complexes; whiskers are 95% bootstrap intervals.

3.10 Effect of similarity to pre-cutoff structures

Accuracy rises weakly but significantly with similarity to the pre-cutoff structural record. Pooled over 245 predictions, the strongest single predictor is not any individual chain but the composite mean identity over the closest *complete* pre-cutoff complex (Spearman $\rho = +0.271$, $p = 1.8 \times 10^{-5}$); the individual components are weaker (TCR β +0.192, peptide +0.151, TCR α +0.127, MHC +0.059, n.s.). The full pre-specified model (DockQ $\sim$ method + TCR α + TCR β + peptide + MHC + class) explains 7.2% of the variance (adjusted $R^2 = 0.048$), with TCR β identity worth +0.041 DockQ per 10 identity points ($p = 0.024$) and peptide identity +0.015 ($p = 0.005$). Between the extreme bins of the composite score, mean DockQ rises from 0.515 to 0.702 (AlphaFold 3) and 0.438 to 0.661 (Protenix).

The negative MHC coefficient in that model is an artefact of no variance rather than evidence that MHC familiarity hurts: MHC identity has median 100% and IQR 99.4–100%, because 111 of 126 targets are class I HLA alleles whose peptide-binding domain is in the PDB many times over. A useful counterexample: two targets have a 100%-identical complete pre-cutoff complex — 7RTR (identical to 7N1F) and 9NDS (identical to 4FTV) — and AlphaFold 3 still scores 7RTR at only 0.36, CAPRI acceptable. Sequence-level availability of the answer does not guarantee a good prediction.

The threshold sweep matters for how method comparisons should be read. Filtering at $\leq 95\%$ maximum TCR identity retains 69 targets and barely moves absolute accuracy (AlphaFold 3 mean DockQ $0.585 \rightarrow 0.547$; Protenix $0.515 \rightarrow 0.506$), but it dissolves the AlphaFold 3–over–Protenix result: the paired advantage falls from $+0.063$ ($p = 0.031$) to $+0.027$ ($p = 0.40$), and at $\leq 90\%$ it reverses sign. Below 95% the retained samples (14, 7 and 1 targets) are too small to conclude anything either way. We therefore conclude that the AlphaFold 3 advantage over Protenix on this benchmark is carried by targets with a near-duplicate TCR in the pre-cutoff PDB. It is also the reason we did not build the primary benchmark on a TCR-identity filter: only 7 of 126 targets fall below 70% identity, and those 7 are not “hard targets” but mostly unusual or non-human TCRs.

3.11 Structural divergence between AlphaFold 3 and ESMFold2

Equal accuracy is not the same as equal predictions. Across the 126 complexes the median C α RMSD between the AlphaFold 3 and ESMFold2 predictions of the same FASTA is 1.35 Å, but the distribution runs to 13.9 Å, and the disagreement is anatomically specific. The two models agree closely on everything except where the TCR sits: median peptide RMSD in the MHC frame is 0.44 Å and median peptide conformation RMSD is 0.33 Å, and the median internal RMSD of the TCR module itself is 0.74 Å — but the median MHC–TCR α and MHC–TCR β interface RMSDs are 1.62 and 1.63 Å, and the rigid-body decomposition gives a median TCR rotation of 11.1° with a tail reaching 177°. On 20 of 126 complexes the two models place the TCR more than 30° apart, and on 10 of them more than 90° apart, while building the same pMHC and the same TCR fold. The disagreement is a docking disagreement, not a folding one.

The disagreement is also bimodal rather than graded. Of the 126 complexes, 106 place the TCR within 30° of each other and only 12 exceed 60°; there is essentially nothing in between, and those 12 are exactly the 12 largest cross-model RMSDs. Superposing each prediction’s MHC on the native’s and measuring its TCR against the native’s separates the two models on these cases: among the 20 complexes that disagree by more than 30°, AlphaFold 3 is the misplaced one on 12, ESMFold2 on 4, and both are misdocked on 4. Over the 10 largest disagreements the median Global DockQ is 0.335 for AlphaFold 3 and 0.581 for ESMFold2, and only 2 of the 10 have both models at medium-or-better accuracy, against 101 of the remaining 116. Where the two architectures disagree about docking, at least one of them has failed, and it is usually AlphaFold 3.

Two geometrically distinct failures produce the disagreement. In the first, the misplaced TCR is rotated 100–177° and translated 50–63 Å: it is not docked on the groove at all but on another surface of the MHC, still burying 55–150 C α contacts there while making no contact with the peptide (9EJG, 8RYO, 7RK7, 7PB2; the fifth such case, 9EJH, is the inverse — its native TCR is the one docked away from the peptide, and AlphaFold 3 builds the canonical peptide-centred pose instead). In the second, the TCR still straddles the groove and still contacts the peptide, but is turned around on it: 70–177° of rotation with the centre of mass only 4–12 Å from where the other

model puts it, a reversal of docking polarity rather than a relocation (9RU5, 8ES9, 7RM4, 7RRG, 9NW2).

The relocations have a specific and avoidable cause. Projecting the TCR centre on the axis from the MHC centre to the peptide centre places the native TCR 28.8 Å above the peptide-binding platform in the median complex, and no class I native below 14 Å. Three predictions — 8RYO, 7RK7 and 7PB2, all AlphaFold 3, all class I — put the TCR 24–26 Å *below* the platform, docked on the β -sheet floor that in the intact molecule is covered by the $\alpha 3$ domain and $\beta 2$ -microglobulin and that the domain truncation of Section 2.2 leaves exposed. On these three complexes AlphaFold 3 scores 0.28, 0.24 and 0.31 while ESMFold2, given the same truncated input, scores 0.91, 0.53 and 0.63 and places the TCR within 3–19° of the native. Part of AlphaFold 3’s tail of docking failures is therefore an artefact of presenting it with a non-physiological surface, not a failure to recognise the correct one — a caveat that applies to any truncated-input benchmark, and one ESMFold2 appears not to share.

That disagreement is a strong, entirely reference-free predictor of accuracy (Figure 6A). Cross-model C α RMSD correlates with AlphaFold 3’s Global DockQ at Spearman $\rho = -0.756$ and with ESMFold2’s at $\rho = -0.670$, and with the *worse* of the two at $\rho = -0.884$ ($p = 1.1 \times 10^{-42}$). For comparison, on the same 126 complexes AlphaFold 3’s own ipTM predicts that worse-of-two value at $\rho = +0.604$ and ESMFold2’s at $\rho = +0.639$. Cross-architecture agreement is therefore a better native-free quality signal than either model’s internal confidence, than pDockQ, and than the composite and random-forest scores of Section 3.9 — and it requires no confidence head, no training, and no calibration.

By divergence quartile (Figure 6B), median Global DockQ falls from 0.860 (AlphaFold 3) and 0.862 (ESMFold2) in the most-agreeing quartile to 0.518 and 0.564 in the most-divergent one. Notably, the two models mostly fail on the *same* complexes: both quartile trends are parallel, and among the 20 most divergent complexes, 4 are failed by both at the medium threshold while 12 are solved by exactly one.

The residual complementarity is nonetheless substantial (Figure 6C). At DockQ ≥ 0.49 , AlphaFold 3 succeeds on 107 of 126 complexes and ESMFold2 on 115; the two together succeed on 119 (103 both, 4 AlphaFold 3 only, 12 ESMFold2 only, 7 neither). Pooling an MSA-based and a language-model-based predictor therefore raises the medium-or-better rate from 85% or 91% to 94% — echoing the pooling gains Yin et al. report for antibody–peptide complexes [13] — and, because disagreement itself flags the failures, the pool comes with a built-in triage signal.

Two related class II entries illustrate the pattern concretely. 7SG1 and 7SG2 are HLA-DQ2 complexes with closely related gluten-derived peptides. On 7SG2 both models are correct (AlphaFold 3 0.706, ESMFold2 0.642) and their predictions differ by 2.72 Å with a 14° TCR rotation. On 7SG1 the models diverge by 3.56 Å with a 26° rotation, and AlphaFold 3 drops to 0.396 while ESMFold2 holds at 0.736. Both models report high confidence on both complexes (AlphaFold 3 ipTM 0.87 and 0.92; ESMFold2 0.81 and 0.83) — the confidence heads do not see the failure, and the disagreement does.

3.12 Denoising trajectories of the two samplers

Instrumenting both diffusion samplers on the same complexes shows that the equality of outcome conceals very different dynamics (Figure 7A). AlphaFold 3 begins denoising from an initial noise scale of $\sim 2,560$ Å and is still ~ 44 Å from the native at the midpoint of its 200-step schedule;




Figure 6: Structural divergence between the AlphaFold 3 and ESMFold2 predictions. (A) Cα RMSD between the AlphaFold 3 and ESMFold2 predictions of the same sequence against the worse of their two Global DockQ scores. (B) Global DockQ by quartile of that divergence; the two methods degrade together. (C) Complexes reaching medium-or-better accuracy: individually, jointly, and neither.

ESMFold2 begins at ~ 256 Å and is already within ~ 5.7 Å at its own midpoint. Measured as the fraction of the schedule at which the trajectory comes permanently within 1 Å of its own final frame, AlphaFold 3 converges at 73% of its steps and ESMFold2 at 58%. AlphaFold 3 spends most of its schedule in a regime where, at the scale the noise schedule sets, no structure has settled at all; ESMFold2 commits early and refines.

The same behaviour at scale (Figure 7B, 30 complexes $\times$ 5 samples = 30,000 recorded frames) puts AlphaFold 3’s convergence in a narrow window. Mean Cα RMSD to the final frame of the same sample falls below 2.0 Å at step 150 and below 1.0 Å at step 158; every complex is within 1.0 Å by step 159; accuracy against the native plateaus at 3.37 Å around step 163; and mean RMSD to the final frame falls below 0.5 Å at step 164. Four independent criteria land between steps 150 and 164 of 200. For roughly the first 140 steps the state is noise at the scale the schedule sets — AlphaFold 3’s schedule holds σ above 5 Å until $\sim 70\%$ of the way through.

3.13 Effect of AlphaFold 3 sampler settings on accuracy and runtime

If the trajectory has effectively converged by step ~ 160 , the last 40 steps and the four extra diffusion samples should be nearly free to remove. They are (Figure 7C, Table 4). Over the same 30 complexes, same MSAs, same seeds, run serially on one idle GPU:

Table 4: **AlphaFold 3 sampler economy on 30 complexes.** Same complexes, same MSAs and templates, same per-complex seed, run serially on one idle RTX PRO 6000. Confidence intervals are paired bootstrap; p values are paired tests against the baseline track. GPU seconds are mean $\pm$ SD over the 30 complexes and exclude a constant 3.3 s of CPU featurisation.

Track	Steps $\times$ samples	GPU s per prediction	Speed-up	Mean Global DockQ	Δ vs baseline (95% CI)	Paired p
Baseline	200 $\times$ 5	10.04 $\pm$ 0.03	1.00 $\times$	0.7212	—	—
Fewer steps	160 $\times$ 5	8.79 $\pm$ 0.03	1.14 $\times$	0.7143	-0.007 (-0.026, +0.007)	0.43
One sample	200 $\times$ 1	6.02 $\pm$ 0.01	1.67 $\times$	0.7143	-0.007 (-0.021, +0.006)	0.34
Both	160 $\times$ 1	5.57 $\pm$ 0.01	1.80$\times$	0.7109	-0.010 (-0.032, +0.007)	0.33

Every confidence interval straddles zero, no track loses a complex from the acceptable band (30/30 everywhere), and the high-accuracy fraction moves from 0.367 to 0.333 — one complex. Adding the constant 3.3 s of CPU featurisation gives the end-to-end figure: 13.4 s $\rightarrow$ 8.9 s per complex, or 401 s $\rightarrow$ 267 s for all 30. Note that the 160-step track is not the 200-step trajectory stopped early — AlphaFold 3 re-discretises the same σ range into 160 coarser steps — so this is an experiment testing the convergence analysis, not a restatement of it.

The speed-up saturates at 1.8 $\times$ for a structural reason. Fitting $t = \text{trunk} + c \times (\text{steps} \times \text{samples})$ across all 120 timed predictions gives a fixed trunk of 4.86 s and 5.08 ms per sample-step ($R^2 = 0.995$). Half of a default AlphaFold 3 run on a complex of this size is the Pairformer trunk and its 10 recycles, which neither shortcut touches; driving diffusion cost to zero would give at most 1.9 $\times$. Dropping from five samples to one also costs the ranking step — AlphaFold 3 normally returns the best of five by ranking score — and on this sample that selection is worth -0.007 DockQ, which is not detectable at $n = 30$.

4 Discussion

4.1 Multiple sequence alignments are not required for TCR-pMHC modelling

The clearest result reported here is a negative one: on 126 post-cutoff complexes, a diffusion predictor conditioned only on a protein language model matched AlphaFold 3 running with server-supplied MSAs and templates (median Global DockQ 0.736 vs 0.751, paired $p = 0.68$), and exceeded it in the fraction of medium-or-better models (91.3% vs 84.9%). This differs from the current consensus [10,11]. The discrepancy appears to be generational rather than contradictory: the PLM predictors in those comparisons are ESMFold-generation models whose structure module descends from AlphaFold 2, whereas ESMFold2 pairs a 6-billion-parameter language model with a diffusion structure head. The most informative data point in Lu et al. is that their best-performing PLM method was the one built on an AlphaFold-3-style architecture [10]; the present result is the same




Figure 7: Denoising trajectories and AlphaFold 3 sampler economy. (A) $C\alpha$ RMSD to the native along the denoising trajectory for four complexes traced in both models; AlphaFold 3 (200 steps, initial $\sigma \approx 2,560$ Å) is still tens of ångströms out at its midpoint, ESMFold2 (68 recorded steps, initial $\sigma \approx 256$ Å) is within a few ångströms. (B) AlphaFold 3 convergence to its own final frame over 30 complexes $\times$ 5 samples; band spans median to 90th percentile. (C) GPU time against mean Global DockQ for four sampler configurations on the same 30 complexes.

observation with the architecture gap closed further. Why this should hold for this problem in particular is worth stating explicitly: TCR variable domains and MHC platforms are among the most heavily sequenced and most heavily crystallised folds in existence, so an MSA adds little that a large language model has not already internalised, whereas the quantity an MSA cannot supply — the relative orientation of two domains with no coevolutionary history with one another — is precisely what remains unsolved.

4.2 Docking, not loop conformation, is the unsolved problem

The per-interface decomposition and the two CDR metrics point the same way. MHC–peptide interfaces are modelled at DockQ ≈ 0.94 by all four methods; CDR3 loops are built to within 1.6–2.3 Å heavy-atom RMSD of the native by all four after framework alignment, with a spread across methods of half an ångström. Docking accuracy, by contrast, spans tens of ångströms, is right-skewed with catastrophic outliers, and accounts for essentially all of the difference between methods and between complexes. This reproduces the decomposition Shi et al. reported [12] on a larger and more recent set, and it argues that further progress on TCR–pMHC modelling will come from the placement problem, not from loop refinement.

4.3 Cross-architecture agreement as a reference-free quality signal

The $C\alpha$ RMSD between an AlphaFold 3 and an ESMFold2 prediction of the same sequence predicted the worse of their two DockQ scores at $\rho = -0.884$, against +0.60 to +0.64 for the models’ own ipTM values, ~ 0.68 for pDockQ, and ~ 0.67 –0.73 for learned quality scores. This is not a subtle effect and it is cheap: one extra prediction from a method that runs in five seconds. It also has a mechanistic reading. Two models that share a generative head but nothing on the conditioning side have largely independent failure modes for the docking problem, so their agreement is evidence in a way that a single model’s self-reported confidence — which, as 7SG1 shows, can be high on a wrong answer — is not. A practical consequence follows: for TCR–pMHC modelling campaigns,

running both architectures, using their agreement as a triage score, and taking the union of their predictions raises the medium-or-better rate on this benchmark from 85% (AlphaFold 3 alone) to 94%.

4.4 Sampler settings and computational cost

AlphaFold 3’s denoising trajectory has converged to within 1 Å of its own answer by step 158 of 200, and running with 160 steps and a single diffusion sample returned the answer in 55% of the GPU time with no detectable accuracy cost over 30 complexes. The scope of that result should be emphasised: $n = 30$, one GPU, one seed per complex, and one narrow size regime (397–416 residues, all in AlphaFold 3’s 512-token bucket). Larger complexes shift the trunk/diffusion split — the trunk scales roughly quadratically in the pair track while diffusion scales with atom count — so the $1.9\times$ ceiling moves. Within this regime, though, the finding is straightforward: for high-throughput TCR–pMHC screening the default settings buy almost nothing, and the compute is better spent on a second architecture, which Section 3.11 shows is worth considerably more than a fifth diffusion sample.

4.5 Implications for benchmark design

Two methodological points generalise beyond this set. First, native construction is not a solved preprocessing step. Forty-six of 136 entries are built from files holding several complete complexes, and the deposited asymmetric unit is often not a complex; six natives required a biological-assembly fetch or a periodic-boundary repair, and one (9J4S) had no physical TCR–pMHC interface in any author-defined assembly, its apparent interface being a lattice contact between two copies. A contact-count check cannot catch this — only requiring all chains to come from a single author-defined assembly does. A benchmark that scores such an entry is scoring noise. Second, redundancy filtering on TCR identity is less informative than it appears for this class of complex: germline V-gene framework dominates a ~ 110 -residue variable domain, so 119 of 126 targets sit above 70% identity to something pre-cutoff, and the $\leq 90\%$ subset is 14 mostly non-human or atypical TCRs rather than a representative hard set. The composite identity to the closest *complete* pre-cutoff complex separates the benchmark far better and at usable sample sizes (46 targets have no complete pre-cutoff match at all). We report both, and we report explicitly that the AlphaFold 3–over–Protenix result does not survive a 95% TCR-identity filter, which is the kind of statement benchmark papers should make routinely.

4.6 Limitations

The class II subset ($n = 15$) is too small to resolve class effects, and TCRmodel2 was run on class I only because it accepts a fixed class II peptide length. Nine entries were excluded for non-standard peptide chemistry or covalent tethering, both of which are biologically interesting categories that sequence-input prediction currently cannot address, so their exclusion reflects a limitation of the methods as much as of the benchmark. “Training set” is a proxy throughout — neither AlphaFold 3’s nor Protenix’s actual training list is public, and we use PDB availability on or before 30 September 2021 as the closest reproducible stand-in, which cannot account for sequence databases or non-TCR structures the models also saw. The reference set for that analysis is only as complete as $\text{TCR3d} \cup \text{STCRDab}$. AlphaFold 3 and Protenix were run through public

servers, so their exact data pipelines are outside our control; ESMFold2 was likewise run through the ESM SDK API, and the trajectory analysis of Section 3.12 used the local biohub/ESMFold2 checkpoint rather than the API model, so its per-frame DockQ values are not interchangeable with the benchmark’s. The composite confidence score of Section 3.9 was developed on this benchmark and requires independent validation. Finally, the sampler-economy result is a single-seed, single-GPU, 30-complex experiment in one size regime, as stated above.

5 Conclusions

On 126 TCR–pMHC complexes released after the shared 30 September 2021 training cutoff, a single-sequence protein-language-model diffusion predictor matched AlphaFold 3 and beat an AlphaFold 3 reproduction, with no MSA and no structural template. Both, and a TCR-specialised AlphaFold 2 derivative, build the pMHC and the CDR3 loops to comparable accuracy and differ almost entirely in where they place the TCR. The disagreement between the two architectures is anatomically confined to that placement, and predicts accuracy better than any of the confidence scores evaluated here, so that a two-architecture ensemble is both more accurate (94% versus 85% medium-or-better) and self-triaging. AlphaFold 3’s denoising trajectory converges by step 158 of 200, and reducing the sampler to 160 steps and one diffusion sample returns the same answers in 55% of the GPU time. Taken together, these observations favour running two architectures at reduced sampler settings over a single architecture at default settings, and using the agreement between them to identify which predictions can be relied upon.

Data and code availability

The benchmark inputs (136 truncated FASTAs, chain manifests, review records), the rebuilt experimental references and their provenance manifest, all predictions, and the complete scoring, analysis and figure code are contained in the accompanying repository. Per-interface DockQ results are provided in long format (one row per interface plus Global DockQ), TCR CDR RMSD results per prediction with full coverage and quality-control columns, CDR IDDT and pLDDT in long format (one row per prediction and region, both contact sets), the sampler-economy tracks with per-complex timings and paired differences, and the per-structure confidence features and bootstrap intervals behind the composite score of Section 2.10. All analyses are reproducible from the recorded commands; software versions and model checkpoint revisions and hashes are recorded in the accompanying manifests.

References

1. Gowthaman R, Pierce BG. TCR3d: The T cell receptor structural repertoire database. *Bioinformatics* 2019;35(24):5323–5325. doi:10.1093/bioinformatics/btz517
2. Leem J, de Oliveira SHP, Krawczyk K, Deane CM. STCRDab: the structural T-cell receptor database. *Nucleic Acids Res* 2018;46(D1):D406–D412. doi:10.1093/nar/gkx971
3. Lin V, Cheung M, Gowthaman R, Eisenberg M, Baker BM, Pierce BG. TCR3d 2.0: expanding the T cell receptor structure database with new structures, tools and interactions. *Nucleic Acids Res* 2025;53(D1):D604–D608. doi:10.1093/nar/gkae840

4. Jumper J, Evans R, Pritzel A, et al. Highly accurate protein structure prediction with AlphaFold. *Nature* 2021;596(7873):583–589. doi:10.1038/s41586-021-03819-2
5. Yin R, Feng BY, Varshney A, Pierce BG. Benchmarking AlphaFold for protein complex modeling reveals accuracy determinants. *Protein Sci* 2022;31(8):e4379. doi:10.1002/pro.4379
6. Abramson J, Adler J, Dunger J, et al. Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature* 2024;630(8016):493–500. doi:10.1038/s41586-024-07487-w
7. Bradley P. Structure-based prediction of T cell receptor:peptide-MHC interactions. *eLife* 2023;12:e82813. doi:10.7554/eLife.82813
8. Yin R, Ribeiro-Filho HV, Lin V, Gowthaman R, Cheung M, Pierce BG. TCRmodel2: high-resolution modeling of T cell receptor recognition using deep learning. *Nucleic Acids Res* 2023;51(W1):W569–W576. doi:10.1093/nar/gkad356
9. Lin Z, Akin H, Rao R, et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science* 2023;379(6637):1123–1130. doi:10.1126/science.ade2574
10. Lu J, Zhu X, Hu X, Zhang C, Feng F. Benchmarking TCR-pMHC structure prediction: a unified evaluation and CDR3-based functional insights. *Brief Bioinform* 2026;27(3):bbag289. doi:10.1093/bib/bbag289
11. Asuncion-París A, Romero-Durana M, Valencia A, Farriol-Duran R, Guallar V. Structural quality-tier assessment for TCR-pMHC functional enrichment. *Front Immunol* 2026;17:1869810. doi:10.3389/fimmu.2026.1869810
12. Shi Y, Parks JM, Smith JC. Comparative analysis of TCR and TCR-pMHC complex structure prediction tools. *J Chem Inf Model* 2025;65(13):7156–7173. doi:10.1021/acs.jcim.5c00298
13. Yin R, Saravanakumar S, Shi SY, et al. Benchmarking AlphaFold and related deep learning approaches for modeling antibody and TCR antigen recognition. *bioRxiv* 2026:2026.07.04.736425. doi:10.64898/2026.07.04.736425
14. Basu S, Wallner B. DockQ: a quality measure for protein-protein docking models. *PLoS One* 2016;11(8):e0161879. doi:10.1371/journal.pone.0161879
15. Mirabello C, Wallner B. DockQ v2: improved automatic quality measure for protein multimers, nucleic acids, and small molecules. *Bioinformatics* 2024;40(10):btae586. doi:10.1093/bioinformatics/btae586
16. Lensink MF, Brysbaert G, Raouraoua N, et al. Impact of AlphaFold on structure prediction of protein complexes: the CASP15-CAPRI experiment. *Proteins* 2023;91(12):1658–1683. doi:10.1002/prot.26609
17. Dunbar J, Deane CM. ANARCI: antigen receptor numbering and receptor classification. *Bioinformatics* 2016;32(2):298–300. doi:10.1093/bioinformatics/btv552
18. Lefranc MP, Pommié C, Ruiz M, et al. IMGT unique numbering for immunoglobulin and T cell receptor variable domains and Ig superfamily V-like domains. *Dev Comp Immunol* 2003;27(1):55–77. doi:10.1016/s0145-305x(02)00039-3
19. Woods LJ, Neff B, Lahouel K, et al. General prediction of T cell receptor antigen specificity from sequence using AlphaFold 3. *bioRxiv* 2026:2026.06.02.729478. doi:10.64898/2026.06.02.729478
20. Bryant P, Pozzati G, Elofsson A. Improved prediction of protein-protein interactions using AlphaFold2. *Nat Commun* 2022;13(1):1265. doi:10.1038/s41467-022-28865-w
21. ESMFold2 model release, biohub/ESMFold2 (revision 8fc3ff47) with biohub/ESMC-6B (revision 45b0fa5d), accessed via the ESM SDK, model esmfold2-fast-2026-05.
22. Protenix: an open reproduction of the AlphaFold 3 architecture (ByteDance AML AI4Science), run through the public Protenix server.
23. Mariani V, Biasini M, Barbato A, Schwede T. IDDT: a local superposition-free score for comparing protein structures and models using distance difference tests. *Bioinformatics*

2013;29(21):2722–2728. doi:10.1093/bioinformatics/btt473